

# **Nishauri 2.0 - Leveraging Retrieval-Augmented Generation for Digital Health**

## **Introduction:**

The challenge posed in the Nishauri application development revolves around integrating Retrieval-Augmented Generation (RAG) to create an advanced chatbot system. This initiative aims to harness Generative AI to enhance digital health services, ensuring accurate and timely information dissemination to users/patients.

## **Understanding of the Case Study:**

In the realm of healthcare, ensuring the precision of information shared is of utmost importance. Here, the integration of Retrieval-Augmented Generation (RAG) models emerges as a crucial strategy, acting as a bridge between dynamic technological advancements and the essential need for accurate healthcare information dissemination.

## **Definition of RAG:**

RAG is a sophisticated model combining a parametric (pre-trained large language model) and a non-parametric model. This fusion enables dynamic information retrieval and generation, ensuring the dissemination of accurate information to the users of Nishauri applications

## **Components of RAG:**

### **1. The Parametric Model:**

Utilizes pre-trained large language models, ensuring robustness and efficiency in generating responses.

Drawbacks and Advantages:

While parametric models excel in their predictive capabilities, they are constrained by resource-intensive fine-tuning processes. However, their advantages lie in their adaptability and effectiveness in generating responses.

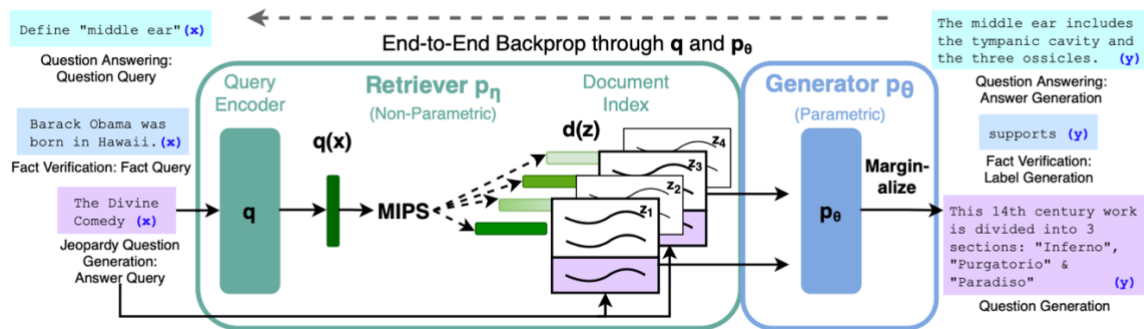
### **2. The Non-parametric Model:**

Employs a dynamic retrieval mechanism for acquiring relevant information.

Drawbacks and Advantages:

Non-parametric models offer flexibility in retrieving information but may lack the predictive capabilities of parametric models. However, they excel in scalability and adaptability.

## **Architecture of the Approach we intend to use**



Overview of the approach step by step :

We combine a pre-trained retriever (*Query Encoder* + *Document Index*) with a pre-trained seq2seq model (*Generator*) and fine-tune end-to-end. For query  $x$ , we use Maximum Inner Product Search (MIPS) to find the top-K documents  $z_i$ . For final prediction  $y$ , we treat  $z$  as a latent variable and marginalize over seq2seq predictions given different documents.

#### 1. Combine Pre-trained Retriever and Seq2Seq Model:

- Integrate a pre-trained retriever, comprising a Query Encoder and Document Index, with a pre-trained seq2seq model, known as the Generator.

#### 2. Fine-tune End-to-End:

- Implement end-to-end fine-tuning of the combined model to ensure seamless interaction between the retriever and the seq2seq model.

#### 3. Query Processing with MIPS:

- When presented with a query (denoted as  $x$ ), employ Maximum Inner Product Search (MIPS) to retrieve the top-K relevant documents (denoted as  $z_i$ ).

#### 4. Prediction Generation:

- Treat the retrieved documents ( $z_i$ ) as latent variables.
- Marginalize over seq2seq predictions considering the varied documents.
- Generate the final prediction (denoted as  $y$ ) based on the marginalized seq2seq predictions.

### Challenges Encountered:

- GPU Requirement:** Fine-tuning pre-trained large language models necessitates significant GPU resources.
- Cost Implications:** Utilizing state-of-the-art parametric models may incur high API usage costs.

3. Memory Constraints: Effective fine-tuning mandates systems with larger memory pools.

**Solutions Implemented:**

1. Google Colab Premium: Subscribing to the premium version of Google Colab addresses both memory and GPU constraints efficiently.
2. Trial Period Access: Exploring access options to pre-trained models like Vertex, albeit temporary, offers potential solutions, although not optimal.